Genaral Questions of DBMS

1. What is a Database Management System (DBMS)?

A **Database Management System (DBMS)** is a software system that enables the creation, management, and manipulation of databases. It provides an interface for users and applications to interact with the data, ensuring data integrity, security, and efficient access.

Key Features:

- **Data Storage**: DBMS stores data in an organized manner.

- **Data Security**: Ensures data confidentiality through user permissions.

- **Data Integrity**: Maintains accuracy and consistency of data.

- **Concurrency Control**: Manages simultaneous access by multiple users.

- **Backup and Recovery**: Provides mechanisms to protect data and restore it in case of failure.

---

2. What are the types of DBMS?

There are four main types of DBMS:

1. **Hierarchical DBMS**: Data is organized in a tree-like structure with a parent-child relationship.

    o Example: IBM's Information Management System (IMS)

2. **Network DBMS**: Data is represented using a graph where entities can have multiple relationships.

    o Example: Integrated Data Store (IDS)

3. **Relational DBMS (RDBMS)**: Data is organized in tables (relations) with rows and columns.

    o Example: MySQL, Oracle, SQL Server

4. **Object-oriented DBMS**: Data is represented as objects, similar to object-oriented programming concepts.

o   Example: db4o

---

3. What is the difference between a primary key and a foreign key?

- **Primary Key**: A column (or set of columns) in a table that uniquely identifies each row in the table. It cannot have NULL values.

    o   Example: emp_id in the **Employee** table.

- **Foreign Key**: A column (or set of columns) that creates a link between two tables. It refers to the primary key in another table and establishes a relationship between the two.

    o   Example: dept_id in the **Employee** table, which refers to the dept_id in the **Department** table.

---

4. What is normalization in DBMS?

Normalization is the process of organizing the data in a database to minimize redundancy and dependency. The goal is to divide large tables into smaller, manageable tables and link them using relationships.

There are several **normal forms (NF)**, each with specific rules:

- **1st Normal Form (1NF)**: Ensures that each column contains atomic values (no repeating groups).

- **2nd Normal Form (2NF)**: Ensures that each non-key column is fully dependent on the primary key.

- **3rd Normal Form (3NF)**: Ensures that all columns are only dependent on the primary key and no transitive dependencies exist.

---

5. What is a transaction in DBMS?

A **transaction** is a sequence of one or more SQL operations executed as a single unit of work. It is used to ensure data consistency and integrity. A transaction has the **ACID properties**:

1. **Atomicity**: The transaction is fully completed or not executed at all.

2. **Consistency**: The transaction brings the database from one valid state to another.

3. **Isolation**: The transaction operates independently of other transactions.

4. **Durability**: Once a transaction is committed, it is permanent, even if the system crashes.

---

6. What is the difference between TRUNCATE and DELETE?

- **TRUNCATE**:
    - Removes all rows from a table.
    - Does not log individual row deletions.
    - Cannot be rolled back in most databases (depends on the DBMS).
    - Faster than DELETE because it doesn't generate individual row logs.
- **DELETE**:
    - Removes rows based on a condition (can be selective).
    - Logs each row deletion.
    - Can be rolled back if in a transaction.
    - Slower than TRUNCATE because of row-level logging.

---

7. What is an index in DBMS?

An **index** is a data structure that improves the speed of data retrieval operations on a table at the cost of additional space and maintenance overhead. It provides a fast way to look up data without having to scan the entire table.

- **Primary Index**: Built automatically when a primary key is defined.
- **Secondary Index**: Can be created on any column to improve query performance.

---

8. What is a view in DBMS?

A **view** is a virtual table that is derived from one or more base tables. It contains a predefined query and is used to present data in a particular format or to simplify complex queries. Views do not store data themselves but display data from underlying tables.

Benefits:

- Simplifies complex queries.

- Provides an extra layer of security by restricting access to specific columns or rows.

- Can be used to encapsulate business logic.

---

9. What is referential integrity?

Referential integrity ensures that relationships between tables are maintained. It prevents actions that would create orphaned records, i.e., records in a child table that do not have a corresponding record in the parent table. This is enforced using **foreign keys**.

For example, if there is a foreign key constraint between the **Employee** and **Department** tables, deleting a department that employees reference should be restricted or cascade to delete the dependent rows in the **Employee** table.

---

10. What are the differences between INNER JOIN and OUTER JOIN?

- **INNER JOIN**: Returns only the rows that have matching values in both tables.

  - Example:

sql

Copy code

SELECT * FROM Employees e INNER JOIN Departments d ON e.dept_id = d.dept_id;

- **OUTER JOIN**: Returns matching rows in addition to rows from one table that do not have corresponding rows in the other table. There are three types:

  - **LEFT OUTER JOIN**: Returns all rows from the left table, and matching rows from the right table (if any).

  - **RIGHT OUTER JOIN**: Returns all rows from the right table, and matching rows from the left table (if any).

  - **FULL OUTER JOIN**: Returns rows when there is a match in either table.

---

11. What is a deadlock in DBMS?

A **deadlock** occurs when two or more transactions are blocked because each transaction is holding a resource and waiting for the other to release a resource. This creates a circular dependency, preventing any of the transactions from completing.

Solution:

- Use **deadlock detection** to identify deadlocks.
- Implement **transaction timeout** to automatically rollback transactions after a certain period.
- **Transaction ordering** and proper resource management can help avoid deadlocks.

---

12. What is a stored procedure in DBMS?

A **stored procedure** is a set of SQL queries that are stored and executed as a single unit. It can accept input parameters, return output, and include logic such as loops and conditions. Stored procedures improve performance and ensure consistency in database operations.

Advantages:

- Reduce network traffic.
- Promote code reusability.
- Help in maintaining security and integrity by centralizing business logic in the database.

---

13. What is the difference between DDL, DML, and DCL?

- **DDL (Data Definition Language)**: Deals with the structure of the database. It includes commands like CREATE, ALTER, DROP, and TRUNCATE.
- **DML (Data Manipulation Language)**: Deals with the manipulation of data within tables. It includes commands like SELECT, INSERT, UPDATE, and DELETE.
- **DCL (Data Control Language)**: Deals with permissions and access control. It includes commands like GRANT and REVOKE.

---

14. What is normalization?

Normalization is the process of organizing data in a database to eliminate redundancy and dependency by dividing large tables into smaller ones and defining relationships between them. This improves data integrity and efficiency.

---

15. What are the ACID properties of a transaction?

The **ACID** properties ensure the reliability of a transaction:

- **Atomicity**: A transaction is either fully completed or not executed at all.

- **Consistency**: The transaction brings the database from one valid state to another.

- **Isolation**: Transactions are isolated from each other, preventing interference.

- **Durability**: Once a transaction is committed, it is permanent, even in the case of system failure.

Stored Procedures

1. What is a stored procedure in SQL? Explain its purpose and benefits.

A **stored procedure** is a precompiled collection of one or more SQL statements that can be executed as a single unit. It is stored in the database and can be invoked by an application or user.
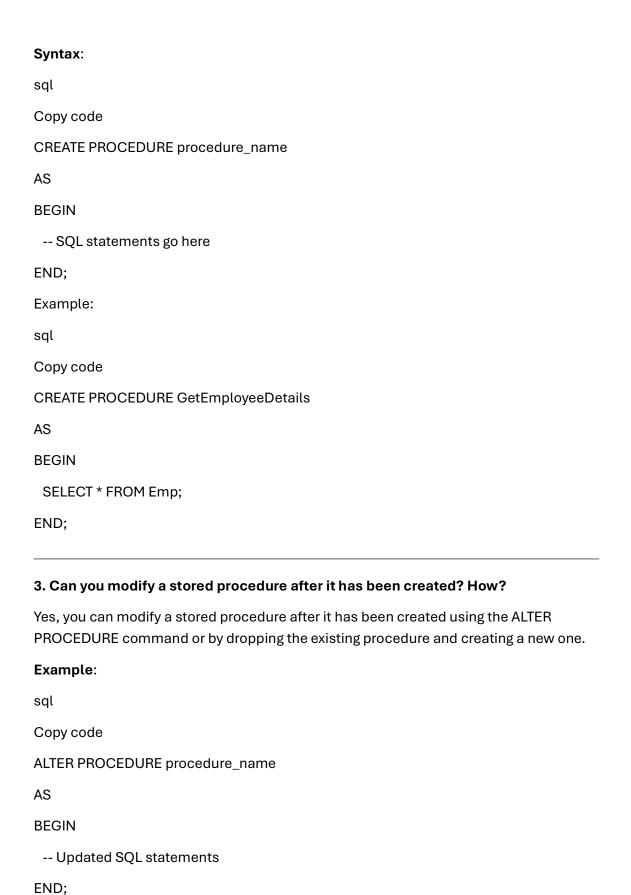
Purpose:

- To simplify repetitive tasks.

- To improve performance by reducing network traffic and compiling SQL queries once.

- To enforce business logic at the database level.

Benefits:

- **Reusability**: You can call stored procedures multiple times.

- **Modularity**: Breaks complex logic into smaller, manageable parts.

- **Security**: Users can execute stored procedures without needing direct access to the underlying data.

- **Performance**: Stored procedures are precompiled, so they execute faster.

---

2. How do you create a stored procedure in SQL? Provide the syntax.

A stored procedure is created using the CREATE PROCEDURE statement followed by the procedure's name and the SQL statements within it.

Syntax:

sql

Copy code

```sql
CREATE PROCEDURE procedure_name

AS

BEGIN

  -- SQL statements go here

END;
```

Example:

sql

Copy code

```sql
CREATE PROCEDURE GetEmployeeDetails

AS

BEGIN

  SELECT * FROM Emp;

END;
```

---

3. Can you modify a stored procedure after it has been created? How?

Yes, you can modify a stored procedure after it has been created using the ALTER PROCEDURE command or by dropping the existing procedure and creating a new one.

Example:

sql

Copy code

```sql
ALTER PROCEDURE procedure_name

AS

BEGIN

  -- Updated SQL statements

END;
```

If ALTER isn't supported, you would drop the procedure and create it again:

sql

Copy code

```sql
DROP PROCEDURE procedure_name;

CREATE PROCEDURE procedure_name

AS

BEGIN

  -- Updated SQL statements

END;
```

---

4. Explain the difference between a stored procedure and a function in SQL.

- **Stored Procedure**: A stored procedure can perform operations like modifying the database (INSERT, UPDATE, DELETE) and doesn't return a value. It may have input and output parameters.

- **Function**: A function always returns a single value (scalar) or a table. It cannot modify database state and is used primarily for computations.

---

5. How can you pass parameters to a stored procedure? Provide an example.

You can pass parameters to a stored procedure during its creation using the IN, OUT, or INOUT keywords, and these parameters can be used in the SQL logic inside the procedure.

Example:

sql

Copy code

```sql
CREATE PROCEDURE GetEmployeeSalary(IN emp_id INT)

AS

BEGIN

  SELECT salary FROM Emp WHERE eno = emp_id;

END;
```

To call the procedure:

sql

Copy code

CALL GetEmployeeSalary(101);

---

6. What is the scope of variables within a stored procedure?

Variables declared inside a stored procedure are local to that procedure. They cannot be accessed outside the procedure, making them private to the procedure's execution.

Example:

sql

Copy code

```sql
CREATE PROCEDURE ExampleProcedure

AS

BEGIN

  DECLARE @Salary INT;

  SET @Salary = 5000;

   -- @Salary can only be used inside this procedure

END;
```

---

7. How do you handle error handling in stored procedures?

You can handle errors in stored procedures using TRY...CATCH blocks. If an error occurs, the CATCH block is executed, allowing for custom error messages or rollback operations.

Example:

sql

Copy code

```sql
BEGIN TRY

  -- SQL statements

  UPDATE Emp SET salary = salary + 100 WHERE eno = 101;
```

END TRY

BEGIN CATCH

   PRINT 'An error occurred';

   -- Optionally, rollback transactions

   ROLLBACK;

END CATCH

---

8. Provide an example where using a stored procedure would be more beneficial than a simple SQL query.

A stored procedure is beneficial when you need to perform complex logic like multiple UPDATE statements, conditional checks, or transactions, and when you need to reuse this logic multiple times.

Example: A procedure to update an employee's salary based on their department and position:

sql

Copy code

```
CREATE PROCEDURE UpdateSalary(IN emp_id INT, IN new_salary DECIMAL)

AS

BEGIN

  IF (SELECT deptno FROM Emp WHERE eno = emp_id) = 10

   BEGIN

     UPDATE Emp SET salary = new_salary WHERE eno = emp_id;

   END

END;
```

This logic can be executed multiple times without writing the query repeatedly.

---

9. Explain the concept of output parameters in stored procedures.

An output parameter allows a stored procedure to return a value to the caller. It can be used to send back results or status information.

Example:

sql

Copy code

```sql
CREATE PROCEDURE GetEmployeeSalary(IN emp_id INT, OUT emp_salary DECIMAL)
AS
BEGIN
  SELECT salary INTO emp_salary FROM Emp WHERE eno = emp_id;
END;
```

When calling the procedure, the output value can be captured:

sql

Copy code

```sql
CALL GetEmployeeSalary(101, @salary);
SELECT @salary;
```

---

10. What is a transaction in a stored procedure, and how do you manage it?

A transaction is a sequence of SQL operations executed as a single unit. It ensures that either all operations are completed successfully or none at all. In stored procedures, you can manage transactions using BEGIN TRANSACTION, COMMIT, and ROLLBACK.

Example:

sql

Copy code

```sql
CREATE PROCEDURE TransferFunds(IN from_acc INT, IN to_acc INT, IN amount DECIMAL)
AS
BEGIN
  BEGIN TRANSACTION;

  UPDATE Accounts SET balance = balance - amount WHERE account_id = from_acc;
  UPDATE Accounts SET balance = balance + amount WHERE account_id = to_acc;
```

```
  IF (/* Check for errors */)

  BEGIN

    ROLLBACK;

  END

  ELSE

  BEGIN

    COMMIT;

  END

END;
```

Functions

1. What is a function in SQL? Explain its purpose and benefits.

A **function** in SQL is a stored program that can accept inputs, perform operations on those inputs, and return a single value or a set of values (in the case of table-valued functions). Functions are used to encapsulate logic that can be reused in SQL queries, improving readability and maintainability.

Purpose:

- To simplify complex logic and encapsulate it in a reusable way.
- To perform operations on data and return results as part of a query.
- To improve performance by centralizing computations.

Benefits:

- **Reusability**: Once defined, functions can be used multiple times.
- **Modularity**: They break complex operations into smaller, manageable pieces.
- **Flexibility**: Functions can be used in SELECT, WHERE, HAVING, or ORDER BY clauses.
- **Code Centralization**: Business logic can be centralized, which helps in maintaining consistency.

---

2. How do you create a function in SQL? Provide the syntax.

A function is created using the CREATE FUNCTION statement, followed by its name, parameters (if any), return type, and the SQL logic.

Syntax:

sql

Copy code

CREATE FUNCTION function_name (parameter1 datatype, parameter2 datatype, ...)

RETURNS return_type

AS

BEGIN

  -- SQL statements

  RETURN value;

END;

Example:

sql

Copy code

CREATE FUNCTION GetEmployeeSalary(emp_id INT)

RETURNS DECIMAL

AS

BEGIN

  DECLARE salary DECIMAL;

  SELECT salary INTO salary FROM Employee WHERE emp_id = emp_id;

  RETURN salary;

END;

This function returns the salary of an employee based on their ID.

---

3. What is the difference between a function and a stored procedure in SQL?

- **Function**:
  - Returns a single value or a table (in the case of table-valued functions).

- o   Cannot perform DML operations like INSERT, UPDATE, or DELETE.

- o   Can be used in a SELECT statement and as part of expressions.

- o   Typically used for computations or data transformation.

- **Stored Procedure**:

  - o   May return multiple result sets and can also perform data manipulation operations.

  - o   Cannot be used in a SELECT statement.

  - o   Can have both input and output parameters.

  - o   Can be used to encapsulate business logic and complex transactions.

---

4. How can you pass parameters to a function in SQL? Provide an example.

You can pass parameters to a function by defining them in the CREATE FUNCTION statement. These parameters can then be used within the function to perform operations.

Example:

sql

Copy code

```
CREATE FUNCTION GetEmployeeSalary(emp_id INT)

RETURNS DECIMAL

AS

BEGIN

   DECLARE salary DECIMAL;

   SELECT salary INTO salary FROM Employee WHERE emp_id = emp_id;

   RETURN salary;

END;
```

To call the function:

sql

Copy code

```
SELECT dbo.GetEmployeeSalary(101);
```

This will return the salary of the employee with emp_id = 101.

---

5. Can a function return a table? If so, how?

Yes, a function can return a table. These types of functions are called **table-valued functions (TVF)**. You define the return type as a table and return a result set.

Syntax:

sql

Copy code

CREATE FUNCTION GetEmployeesByDept(department_id INT)

RETURNS TABLE

AS

RETURN

(

    SELECT emp_id, emp_name, salary FROM Employee WHERE dept_id = department_id

);

To use the function:

sql

Copy code

SELECT * FROM dbo.GetEmployeesByDept(10);

This will return a table of employees who work in department 10.

---

6. What are scalar functions in SQL?

A **scalar function** returns a single value (such as INT, DECIMAL, or VARCHAR) for each row processed. Scalar functions can be used in SELECT statements, WHERE clauses, and other expressions where a single value is required.

Example:

sql

Copy code

CREATE FUNCTION GetFullName(first_name VARCHAR(50), last_name VARCHAR(50))

RETURNS VARCHAR(100)

AS

BEGIN

  RETURN CONCAT(first_name, ' ', last_name);

END;

Usage:

sql

Copy code

SELECT dbo.GetFullName(first_name, last_name) FROM Employee;

This function returns the full name by concatenating the first and last names of the employee.

---

7. Can you perform DML operations inside a function in SQL?

No, you cannot perform DML operations (INSERT, UPDATE, DELETE) inside a function in SQL. Functions are intended for computations or returning data, not for modifying the database state. However, stored procedures can perform DML operations.

---

8. What is the difference between COALESCE and ISNULL in SQL?

- **COALESCE**: Returns the first non-NULL expression from a list of expressions. It can handle more than two parameters.

Example:

sql

Copy code

SELECT COALESCE(NULL, NULL, 'FirstNonNull', 'SecondNonNull');

-- Result: 'FirstNonNull'

- **ISNULL**: Checks if a single expression is NULL and replaces it with another value. It only takes two arguments.

Example:

sql

Copy code

SELECT ISNULL(salary, 0) FROM Employee;

---

9. What are the advantages of using functions in SQL?

- **Reusability**: Functions can be called multiple times in different queries.

- **Modularity**: Functions encapsulate logic, making queries easier to read and maintain.

- **Performance**: Predefined functions can improve query performance by avoiding repetitive calculations.

- **Consistency**: Using functions ensures that the same logic is applied consistently across different queries or applications.

- **Data abstraction**: Functions can abstract complex logic from the users and applications.

---

10. What is a recursive function in SQL?

A **recursive function** is a function that calls itself during its execution. It is useful for operations like hierarchical data retrieval (e.g., organizational structures) or traversing tree-like structures.

However, recursive queries (using recursive CTEs) are more common in SQL than recursive functions.

Example: In SQL, recursive functions are usually implemented with a **Common Table Expression (CTE)**:

sql

Copy code

WITH RECURSIVE EmployeeHierarchy AS

(

   SELECT emp_id, emp_name, manager_id

   FROM Employee

   WHERE manager_id IS NULL

   UNION ALL

SELECT e.emp_id, e.emp_name, e.manager_id

FROM Employee e

INNER JOIN EmployeeHierarchy eh ON e.manager_id = eh.emp_id

)

SELECT * FROM EmployeeHierarchy;

This retrieves all employees and their hierarchical structure based on the manager_id.

VIEW

Let's proceed with the next title: **"Views in SQL"**. Below are detailed answers to the 10 questions about views:

1. What is a View in SQL?

A **view** in SQL is a virtual table that is created by executing a SELECT query on one or more base tables. It does not store data itself; instead, it displays data stored in underlying tables, making complex queries more manageable by abstracting and simplifying the structure.

Benefits of Views:

- **Simplified Querying**: Complex joins and subqueries can be encapsulated into a single view, making querying easier.

- **Security**: Views can restrict access to sensitive data by showing only specific columns or rows.

- **Consistency**: A view ensures that the same query logic is used across multiple places, preventing redundancy.

---

2. How do you create a View in SQL? Provide an example.

You can create a view using the CREATE VIEW statement. The syntax involves defining the view's name, followed by a SELECT query that specifies the columns and tables from which the view will be derived.

Syntax:

sql

Copy code

CREATE VIEW view_name AS

SELECT column1, column2, …

FROM table_name

WHERE condition;

Example:

sql

Copy code

CREATE VIEW EmployeeSalary AS

SELECT emp_id, emp_name, salary

FROM Employee

WHERE salary > 50000;

This view EmployeeSalary will show only employees whose salary is greater than 50,000.

---

3. What are the benefits of using Views in SQL?

Benefits:

- **Abstraction**: Views abstract the complexity of the underlying tables and queries.

- **Security**: Views can limit the exposure of sensitive data by only allowing access to selected columns.

- **Consistency**: Using a view ensures that queries are consistent across different applications and users.

- **Performance**: Views can improve performance by simplifying complex joins or aggregations into reusable components.

---

4. Can we update data through a View in SQL?

Yes, it is possible to update data through a view, but only if the view is updatable. A view is considered updatable if it:

- Refers to a single base table.

- Does not contain aggregate functions (e.g., SUM, AVG).

- Does not contain DISTINCT, GROUP BY, or JOIN clauses.

For example, a simple view that directly maps to a single table can be updated:

sql

Copy code

```
CREATE VIEW EmployeeView AS

SELECT emp_id, emp_name, salary

FROM Employee;


UPDATE EmployeeView

SET salary = 60000

WHERE emp_id = 101;
```

However, if the view involves complex joins or aggregations, it may not be directly updatable.

---

5. What is an Indexed View in SQL?

An **indexed view** is a view that has a unique clustered index created on it. This allows SQL Server (or another DBMS) to store the result of the view physically, improving performance for frequently accessed views.

Syntax:

sql

Copy code

```
CREATE UNIQUE CLUSTERED INDEX idx_name ON view_name (column1);
```

Use Case: Indexed views are used to improve query performance when the view contains complex joins or aggregations that are frequently queried.

---

6. What is a Materialized View in SQL?

A **materialized view** is a view whose result set is stored physically in the database, unlike a regular view which is virtual. Materialized views are typically used to improve performance by storing the results of expensive queries.

Benefits:

- **Performance**: They provide fast access to the results of complex queries by storing precomputed data.

- **Refresh Options**: Materialized views can be set to refresh at specific intervals (manual or automatic).

Example:

sql

Copy code

```
CREATE MATERIALIZED VIEW SalesSummary AS

SELECT product_id, SUM(quantity) AS total_sales

FROM Sales

GROUP BY product_id;
```

The data in this view can be refreshed periodically to keep it up to date.

---

7. How do you drop a View in SQL?

You can drop a view using the DROP VIEW statement.

Syntax:

sql

Copy code

```
DROP VIEW view_name;
```

Example:

sql

Copy code

```
DROP VIEW EmployeeSalary;
```

This removes the view EmployeeSalary from the database.

---

8. Can we use joins in Views?

Yes, views can include joins. You can create a view that joins multiple tables, allowing users to retrieve data from those tables as though it were coming from a single source.

Example:

sql

Copy code

```
CREATE VIEW EmployeeDept AS

SELECT e.emp_id, e.emp_name, d.dept_name

FROM Employee e

JOIN Department d ON e.dept_id = d.dept_id;
```

This view combines employee data and department data through an inner join.

---

9. What is the difference between a View and a Table in SQL?

- **View**: A virtual table created by a SELECT query. It does not store data itself, and its data is retrieved from the underlying base tables whenever queried.

- **Table**: A physical structure that stores data. It is the actual storage location for data in a database.

Differences:

- Views simplify queries and abstract complexity, while tables store actual data.

- A table is persistent, but a view is just a query that is executed each time it is accessed.

- Views can provide a layer of security and allow users to access only specific data.

---

10. How can Views improve query performance in SQL?

Views can improve query performance by:

- **Reducing Complexity**: Complex joins, subqueries, and aggregations can be encapsulated in a view, making the main query simpler and more readable.

- **Encapsulation of Repeated Logic**: Views can encapsulate business logic or calculations that are used frequently, reducing the need to repeat complex SQL logic in every query.
- **Materialized Views**: If using materialized views, they store the query result physically, improving performance for frequently queried data.

Let's move on to the next title: **"Subqueries in SQL"**. Below are detailed answers to the 10 questions about subqueries.

1. What is a Subquery in SQL?

A **subquery** is a query embedded within another query. It is typically used to perform operations that involve filtering, aggregation, or comparison of results that are derived from a second query. Subqueries can be placed in the SELECT, FROM, WHERE, or HAVING clause.

Example:

sql

Copy code

```
SELECT emp_name, salary

FROM Employee

WHERE dept_id = (SELECT dept_id FROM Department WHERE dept_name = 'Sales');
```

This subquery retrieves the dept_id of the 'Sales' department and uses it to filter employees from that department.

---

2. How do you write a Subquery in SQL? Provide an example.

A subquery is written by embedding a query within parentheses inside the main query. It can be used in various clauses like SELECT, WHERE, or FROM.

Syntax:

sql

Copy code

```
SELECT column_name

FROM table_name

WHERE column_name = (SELECT column_name FROM table_name WHERE condition);
```

Example:

sql

Copy code

SELECT emp_name, salary

FROM Employee

WHERE salary > (SELECT AVG(salary) FROM Employee);

This query retrieves the names and salaries of employees whose salary is greater than the average salary.

---

3. What are the types of Subqueries in SQL?

There are mainly three types of subqueries:

1. **Single-row Subquery**: Returns a single row and a single column. It is used with operators like =, >, <, etc.

   o   Example: SELECT * FROM Employee WHERE salary = (SELECT MAX(salary) FROM Employee);

2. **Multiple-row Subquery**: Returns multiple rows. It is used with operators like IN, ANY, ALL.

   o   Example: SELECT * FROM Employee WHERE dept_id IN (SELECT dept_id FROM Department WHERE location = 'New York');

3. **Correlated Subquery**: A subquery that references columns from the outer query. The subquery is evaluated once for each row processed by the outer query.

   o   Example:

sql

Copy code

SELECT emp_name

FROM Employee e

WHERE salary > (SELECT AVG(salary) FROM Employee WHERE dept_id = e.dept_id);

The subquery is evaluated for each employee based on their department.

---

4. What is a Correlated Subquery?

A **correlated subquery** refers to a subquery that depends on the outer query for its values. It uses a column from the outer query in its condition, which means it is executed repeatedly for each row processed by the outer query.

Example:

sql

Copy code

SELECT emp_name

FROM Employee e

WHERE salary > (SELECT AVG(salary) FROM Employee WHERE dept_id = e.dept_id);

Here, for each employee, the subquery calculates the average salary for that employee's department.

---

5. What is the difference between a Subquery and a Join?

- **Subquery**: A subquery is a query nested inside another query. It is often used to retrieve a single value or set of values to compare or filter results. Subqueries can be used in SELECT, WHERE, or FROM clauses.

Example:

sql

Copy code

SELECT emp_name

FROM Employee

WHERE dept_id = (SELECT dept_id FROM Department WHERE dept_name = 'Sales');

- **Join**: A join is used to combine rows from two or more tables based on a related column between them. Joins are typically used when you need to combine data from multiple tables in a single result set.

Example:

sql

Copy code

SELECT e.emp_name, d.dept_name

FROM Employee e

JOIN Department d ON e.dept_id = d.dept_id

WHERE d.dept_name = 'Sales';

Key Difference:

- Subqueries are executed independently, while joins combine multiple tables together directly.

---

6. What is the difference between an Inline View and a Subquery?

- **Inline View**: An inline view is essentially a subquery used in the FROM clause of a query. It is treated like a virtual table within the query and can be used just like any other table.

Example:

sql

Copy code

SELECT emp_name

FROM (SELECT emp_name, salary FROM Employee WHERE salary > 50000) AS HighSalaryEmployees;

- **Subquery**: A subquery can be placed in SELECT, WHERE, or HAVING clauses and is not limited to the FROM clause.

---

7. Can a Subquery return multiple values?

Yes, a subquery can return multiple values, especially when used with operators like IN, ANY, or ALL. These types of subqueries are typically used in WHERE clauses to filter the results based on multiple conditions.

Example:

sql

Copy code

SELECT emp_name

FROM Employee

WHERE dept_id IN (SELECT dept_id FROM Department WHERE location = 'New York');

This subquery returns multiple dept_id values for departments located in 'New York'.

8. What is a Nested Subquery?

A **nested subquery** is a subquery that is placed inside another subquery. A nested subquery allows multiple levels of queries to be performed. Each subquery can depend on the results of the one before it.

Example:

sql

Copy code

SELECT emp_name

FROM Employee

WHERE dept_id = (SELECT dept_id FROM Department WHERE dept_name = (SELECT dept_name FROM Department WHERE location = 'New York'));

Here, the innermost subquery first retrieves the department name for 'New York', which is then used by the second subquery to find the dept_id of that department.

---

9. What are the advantages and disadvantages of using Subqueries?

Advantages:

- **Simplicity**: Subqueries can simplify queries by breaking down complex logic into smaller parts.

- **Reusability**: Subqueries allow you to use the result of one query in another query.

- **Clarity**: They help separate filtering or calculations, making the main query easier to read.

Disadvantages:

- **Performance**: Subqueries can be inefficient, especially if they need to be executed multiple times (as in correlated subqueries).

- **Complexity**: In nested or correlated subqueries, debugging and understanding the logic can become difficult.

---

10. What are the performance considerations when using Subqueries?

- **Execution Time**: A subquery can be executed multiple times, especially if it is correlated, which can significantly impact performance.

- **Indexes**: If the subquery involves a table that is not properly indexed, it can slow down execution.

- **Alternative**: Sometimes, using JOIN instead of subqueries can improve performance, especially when dealing with large datasets.

Let's proceed with the next title: **"Joins in SQL"**. Below are detailed answers to the 10 questions about joins.

Join

1. What is a Join in SQL?

A **join** in SQL is used to combine rows from two or more tables based on a related column between them. Joins help retrieve data that is spread across multiple tables, making it easier to perform complex queries involving multiple entities.

There are different types of joins:

- **INNER JOIN**

- **LEFT JOIN (or LEFT OUTER JOIN)**

- **RIGHT JOIN (or RIGHT OUTER JOIN)**

- **FULL JOIN (or FULL OUTER JOIN)**

- **CROSS JOIN**

- **SELF JOIN**

---

2. What are the different types of Joins in SQL?

1. **INNER JOIN**: This join returns only the rows that have matching values in both tables. It excludes rows with no match.

    o **Example**:

sql

Copy code

```sql
SELECT e.emp_name, d.dept_name
FROM Employee e
INNER JOIN Department d ON e.dept_id = d.dept_id;
```

2. **LEFT JOIN (or LEFT OUTER JOIN)**: This join returns all rows from the left table, and the matching rows from the right table. If there is no match, NULL values are returned for columns from the right table.

   o **Example**:

sql

Copy code

```sql
SELECT e.emp_name, d.dept_name
FROM Employee e
LEFT JOIN Department d ON e.dept_id = d.dept_id;
```

3. **RIGHT JOIN (or RIGHT OUTER JOIN)**: This join is similar to LEFT JOIN, but it returns all rows from the right table and matching rows from the left table.

   o **Example**:

sql

Copy code

```sql
SELECT e.emp_name, d.dept_name
FROM Employee e
RIGHT JOIN Department d ON e.dept_id = d.dept_id;
```

4. **FULL JOIN (or FULL OUTER JOIN)**: This join returns all rows when there is a match in either the left or the right table. If there is no match, NULL values are returned for the non-matching side.

   o **Example**:

sql

Copy code

```sql
SELECT e.emp_name, d.dept_name
FROM Employee e
```

FULL OUTER JOIN Department d ON e.dept_id = d.dept_id;

5. **CROSS JOIN**: This join returns the Cartesian product of the two tables, meaning it combines every row from the first table with every row from the second table.

   o **Example**:

sql

Copy code

SELECT e.emp_name, d.dept_name

FROM Employee e

CROSS JOIN Department d;

6. **SELF JOIN**: A self join is a join where a table is joined with itself. This is useful when you need to compare rows within the same table.

   o **Example**:

sql

Copy code

SELECT e1.emp_name, e2.emp_name

FROM Employee e1

INNER JOIN Employee e2 ON e1.manager_id = e2.emp_id;

---

3. What is the difference between INNER JOIN and OUTER JOIN?

- **INNER JOIN**: Returns only the rows where there is a match in both tables. Rows without a match are excluded.

- **OUTER JOIN**: Returns all rows from one table and matching rows from the other. If there is no match, NULL values are returned for the non-matching rows.

Example:

- **INNER JOIN**: Retrieves employees who are assigned to departments.

- **LEFT JOIN**: Retrieves all employees, including those without a department (with NULL for department details).

---

4. What is the syntax for using a JOIN in SQL?

The basic syntax for using a join involves specifying the type of join, followed by the condition that links the tables together.

Syntax:

sql

Copy code

SELECT column1, column2, …

FROM table1

JOIN table2 ON table1.column = table2.column;

Example (INNER JOIN):

sql

Copy code

SELECT emp_name, dept_name

FROM Employee

INNER JOIN Department ON Employee.dept_id = Department.dept_id;

---

5. What is a Cartesian Product (CROSS JOIN) in SQL?

A **CROSS JOIN** (also known as a Cartesian product) combines every row of the first table with every row of the second table. It is used when you want to pair all possible combinations between two tables.

Example:

sql

Copy code

SELECT emp_name, dept_name

FROM Employee

CROSS JOIN Department;

If Employee has 10 rows and Department has 5 rows, the result will contain 50 rows (10 * 5).

---

6. What is a Self Join in SQL?

A **self join** is when a table is joined with itself. This can be useful for finding relationships between rows within the same table, such as finding employees who work under the same manager.

Example:

sql

Copy code

```
SELECT e1.emp_name, e2.emp_name AS manager_name

FROM Employee e1

INNER JOIN Employee e2 ON e1.manager_id = e2.emp_id;
```

In this case, we are joining the Employee table with itself to find employees and their respective managers.

---

7. How do you optimize performance with Joins in SQL?

Here are some ways to optimize joins:

- **Indexing**: Create indexes on columns that are frequently used in JOIN conditions. This speeds up the matching process.

- **Limit the number of rows**: Use WHERE clauses to filter rows before performing a join.

- **Choose appropriate join types**: Use INNER JOIN when possible to reduce unnecessary rows. Avoid CROSS JOIN unless needed.

- **Avoid joining large tables**: If possible, filter or aggregate data before joining large tables.

---

8. Can you join more than two tables in SQL?

Yes, you can join more than two tables in SQL. You simply continue specifying the join conditions for each additional table.

Example:

sql

Copy code

```
SELECT e.emp_name, d.dept_name, m.manager_name
```

FROM Employee e

INNER JOIN Department d ON e.dept_id = d.dept_id

INNER JOIN Manager m ON e.manager_id = m.manager_id;

Here, we are joining three tables: Employee, Department, and Manager.

---

9. What is the difference between a LEFT JOIN and a RIGHT JOIN?

- **LEFT JOIN**: Returns all rows from the left table and matching rows from the right table. If there is no match, NULL values are returned for the right table.
- **RIGHT JOIN**: Returns all rows from the right table and matching rows from the left table. If there is no match, NULL values are returned for the left table.

Example:

- **LEFT JOIN**: Retrieves all employees, even those without a department.
- **RIGHT JOIN**: Retrieves all departments, even those without employees.

---

10. What is the purpose of JOIN conditions in SQL?

The purpose of **JOIN conditions** is to define how rows from two or more tables should be matched based on common columns. These conditions are usually defined using the ON keyword, specifying a relationship between the tables' columns.

Example:

sql

Copy code

SELECT emp_name, dept_name

FROM Employee

INNER JOIN Department ON Employee.dept_id = Department.dept_id;

Here, the condition Employee.dept_id = Department.dept_id ensures that the employee data is joined with the correct department data.

These answers provide a comprehensive understanding of **joins in SQL**, explaining different types of joins, their syntax, use cases, and how to optimize performance when

using joins. Understanding joins is essential for working with multiple related tables in a relational database.

Let's move on to the next title: **"Triggers in SQL"**. Below are detailed answers to the 10 questions about triggers.

Triggers

1. What is a Trigger in SQL?

A **trigger** is a special type of stored procedure that automatically executes or fires when certain events occur on a specific table or view in a database. These events could include actions like INSERT, UPDATE, or DELETE. Triggers help automate tasks such as enforcing business rules, data validation, or auditing changes.

Example:

sql

Copy code

CREATE TRIGGER before_insert_emp

BEFORE INSERT ON Employee

FOR EACH ROW

SET NEW.salary = 50000;

This trigger automatically sets the salary of any new employee to 50,000 before the insertion into the Employee table.

---

2. What are the types of Triggers in SQL?

There are mainly two types of triggers:

1. **BEFORE Trigger**: Executes before an insert, update, or delete operation. It is often used to modify data before it is committed to the table.

   o **Example**: A BEFORE INSERT trigger that automatically sets default values.

2. **AFTER Trigger**: Executes after an insert, update, or delete operation. It is used to enforce constraints, log changes, or perform additional actions after the data has been modified.

   o **Example**: An AFTER INSERT trigger that logs new entries in an audit table.

---

3. What is the syntax for creating a Trigger in SQL?

The syntax for creating a trigger involves defining the trigger name, timing (BEFORE or AFTER), the event (like INSERT, UPDATE, DELETE), and the body of the trigger.

Syntax:

sql

Copy code

```
CREATE TRIGGER trigger_name

trigger_time trigger_event

ON table_name

FOR EACH ROW

BEGIN

  -- Trigger logic

END;
```

Example:

sql

Copy code

```
CREATE TRIGGER update_salary

AFTER UPDATE ON Employee

FOR EACH ROW

BEGIN

  INSERT INTO Salary_Log (emp_id, old_salary, new_salary, update_time)

  VALUES (OLD.emp_id, OLD.salary, NEW.salary, NOW());

END;
```

This trigger logs changes to the employee salary into the Salary_Log table after an update.

---

4. What are OLD and NEW keywords in Triggers?

The **OLD** and **NEW** keywords are used in triggers to refer to the values of a row before and after an operation:

- **OLD**: Refers to the value before the operation (used in UPDATE or DELETE).

- **NEW**: Refers to the value after the operation (used in INSERT or UPDATE).

Example (for UPDATE):

sql

Copy code

```
CREATE TRIGGER salary_update

AFTER UPDATE ON Employee

FOR EACH ROW

BEGIN

  IF OLD.salary < NEW.salary THEN

    -- Logic for salary increase

  END IF;

END;
```

Here, OLD.salary refers to the employee's salary before the update, and NEW.salary refers to the salary after the update.

---

5. What is the purpose of Triggers in SQL?

Triggers are used to:

- **Enforce data integrity**: Automatically check or modify data before or after an operation.

- **Audit changes**: Automatically track changes to data, such as who updated a record and when.

- **Business logic enforcement**: Automatically apply rules or default values during INSERT or UPDATE.

- **Automation**: Perform operations such as logging, notifications, or sending emails based on certain actions.

---

6. Can a Trigger be used to call another Trigger?

Yes, triggers can **call other triggers** indirectly. For example, one trigger can modify data in such a way that it causes another trigger to fire. However, it's important to avoid creating circular or infinite trigger calls, where one trigger repeatedly triggers another in a loop.

Example:

sql

Copy code

```sql
CREATE TRIGGER before_insert_emp

BEFORE INSERT ON Employee

FOR EACH ROW

BEGIN

  SET NEW.salary = 50000;

END;


CREATE TRIGGER after_insert_emp

AFTER INSERT ON Employee

FOR EACH ROW

BEGIN

  INSERT INTO Audit (emp_id, action) VALUES (NEW.emp_id, 'Inserted');

END;
```

In this case, the after_insert_emp trigger runs after the before_insert_emp trigger inserts a record into the Employee table.

---

7. What is the difference between a Trigger and a Stored Procedure?

- **Trigger**: Automatically executes in response to specific events (like INSERT, UPDATE, or DELETE) on a table or view. Triggers cannot be called explicitly.

- **Stored Procedure**: A precompiled collection of SQL statements that can be executed manually. Stored procedures are invoked explicitly by the user or application.

Key Differences:

- Triggers are event-driven, whereas stored procedures are manually invoked.

- Triggers operate implicitly, while stored procedures are called explicitly using the EXEC command.

---

8. What are the limitations of Triggers in SQL?

- **Performance overhead**: Triggers can slow down operations (like INSERT, UPDATE, or DELETE) because they add additional processing each time these operations occur.

- **Complexity**: Triggers can introduce complexity in the database schema and make it harder to understand the flow of data.

- **No explicit control**: You cannot explicitly call triggers; they run automatically, which can lead to unexpected behavior in some cases.

- **Limited transaction control**: Triggers cannot commit or roll back transactions; they operate within the transaction that caused them.

---

9. Can Triggers be used for data validation?

Yes, triggers can be used to **validate data** before or after an insert, update, or delete operation. For example, you could create a BEFORE INSERT trigger that ensures certain conditions are met, such as checking if an employee's salary is within a valid range.

Example:

sql

Copy code

CREATE TRIGGER check_salary

BEFORE INSERT ON Employee

FOR EACH ROW

BEGIN

  IF NEW.salary < 30000 THEN

SIGNAL SQLSTATE '45000' SET MESSAGE_TEXT = 'Salary must be at least 30000';

   END IF;

END;

In this case, the trigger raises an error if the salary being inserted is below 30,000.

---

10. What are the events that can trigger a Trigger in SQL?

Triggers can be activated by the following events:

1. **INSERT**: Fires after a row is inserted into a table.

2. **UPDATE**: Fires when an existing row in a table is updated.

3. **DELETE**: Fires when a row is deleted from a table.

These events can occur **BEFORE** or **AFTER** the operation, depending on how the trigger is defined.

Example (DELETE):

sql

Copy code

```
CREATE TRIGGER after_delete_employee

AFTER DELETE ON Employee

FOR EACH ROW

BEGIN

  INSERT INTO Audit_Log (action, emp_id, delete_time)

  VALUES ('DELETE', OLD.emp_id, NOW());

END;
```

This trigger logs the deletion of an employee into the Audit_Log table.

These answers provide a detailed overview of **triggers in SQL**, including their usage, types, and best practices. Triggers are powerful tools for automating database tasks, ensuring data integrity, and enforcing business rules within a database.

Cursors

Let's move on to the next title: **"Cursors in SQL"**. Below are detailed answers to the 10 questions about cursors.

1. What is a Cursor in SQL?

A **cursor** in SQL is a database object used to retrieve, manipulate, and navigate through a result set one row at a time. Cursors are primarily used when you need to perform row-by-row operations on a result set (such as processing individual records in a loop), which can't be done efficiently using a single SQL query.

A cursor allows you to:

- Iterate over query results
- Fetch rows
- Update or delete rows in a loop

Example:

sql

Copy code

```
DECLARE cursor_example CURSOR FOR

SELECT emp_name FROM Employee;
```

Here, a cursor is declared to iterate over the emp_name column of the Employee table.

---

2. What are the types of Cursors in SQL?

There are two main types of cursors in SQL:

1. **Implicit Cursors**: These are automatically created by the database system when a query is executed (like SELECT statements). They don't require manual management.

   o **Example**: A simple SELECT query.

2. **Explicit Cursors**: These are explicitly declared by the programmer and give control over the fetching and navigating of rows in the result set. Explicit cursors allow greater control over the process.

   o **Example**: A cursor declared and used in a stored procedure or function.

---

3. What is the syntax to declare a Cursor in SQL?

The basic syntax for declaring a cursor involves defining the cursor and the SQL query it will execute.

Syntax:

sql

Copy code

```
DECLARE cursor_name CURSOR FOR

SELECT column1, column2

FROM table_name

WHERE condition;
```

Example:

sql

Copy code

```
DECLARE employee_cursor CURSOR FOR

SELECT emp_name FROM Employee WHERE dept_id = 1;
```

This cursor selects the names of employees in department 1.

---

4. What are the steps to use a Cursor in SQL?

Using a cursor involves several steps:

1. **Declare the cursor**: Define the cursor and the query that it will fetch data from.
2. **Open the cursor**: Open the cursor to execute the SQL query and retrieve the result set.
3. **Fetch rows**: Retrieve rows one at a time from the cursor.
4. **Close the cursor**: Once the processing is done, close the cursor to release the resources.

Example:

sql

Copy code

```
DECLARE employee_cursor CURSOR FOR
```

SELECT emp_name FROM Employee WHERE dept_id = 1;

OPEN employee_cursor;

FETCH NEXT FROM employee_cursor INTO @emp_name;

-- Processing the data

CLOSE employee_cursor;

DEALLOCATE employee_cursor;

---

5. What is the difference between an Implicit Cursor and an Explicit Cursor?

- **Implicit Cursor**: Automatically created by the database system when executing a single SQL statement. It doesn't require explicit declaration or management by the programmer.
    - Used for simple queries that return one result set.
    - Example: SELECT * FROM Employee;
- **Explicit Cursor**: Declared and controlled by the programmer. It allows for more complex queries, multiple result sets, and fine control over fetching and updating data.
    - Used when row-by-row processing or multiple result sets are needed.
    - Example: Using a DECLARE statement to declare a cursor for processing rows.

---

6. What are the different types of Cursor Operations in SQL?

Here are the main operations that can be performed on cursors:

1. **DECLARE**: Defines the cursor and the SQL query it will run.
2. **OPEN**: Executes the SQL query and allocates resources for the cursor.
3. **FETCH**: Retrieves the next row from the cursor and stores the data into variables.

4. **CLOSE**: Releases the resources held by the cursor.

5. **DEALLOCATE**: Removes the cursor from memory.

Example:

sql

Copy code

DECLARE emp_cursor CURSOR FOR

SELECT emp_name FROM Employee;

OPEN emp_cursor;

FETCH NEXT FROM emp_cursor INTO @emp_name;

-- Perform processing

CLOSE emp_cursor;

DEALLOCATE emp_cursor;

---

7. What is the FETCH statement in SQL Cursor?

The **FETCH** statement is used to retrieve the next row of data from an open cursor and store it into variables. The cursor must be open before a FETCH statement can be executed.

Syntax:

sql

Copy code

FETCH NEXT FROM cursor_name INTO variable1, variable2, …;

Example:

sql

Copy code

FETCH NEXT FROM emp_cursor INTO @emp_name;

This fetches the next row from the cursor into the @emp_name variable.

---

8. What are Cursor Attributes in SQL?

Cursor attributes provide status information about the cursor. These attributes help you control the flow of execution when working with cursors.

- **%FOUND**: Returns TRUE if the last fetch was successful.

- **%NOTFOUND**: Returns TRUE if the last fetch failed (i.e., no more rows).

- **%ROWCOUNT**: Returns the number of rows fetched so far.

- **%ISOPEN**: Returns TRUE if the cursor is open.

Example:

sql

Copy code

```
FETCH NEXT FROM emp_cursor INTO @emp_name;

IF emp_cursor%NOTFOUND THEN

   -- Handle no data found

END IF;
```

---

9. Can you update data using Cursors in SQL?

Yes, you can use a cursor to update data row by row. After fetching each row, you can perform an UPDATE on the current record.

Example:

sql

Copy code

```
DECLARE emp_cursor CURSOR FOR

SELECT emp_id, salary FROM Employee WHERE dept_id = 1;


OPEN emp_cursor;


FETCH NEXT FROM emp_cursor INTO @emp_id, @salary;
```

```
WHILE @@FETCH_STATUS = 0

BEGIN

  UPDATE Employee

  SET salary = salary * 1.1

  WHERE emp_id = @emp_id;


  FETCH NEXT FROM emp_cursor INTO @emp_id, @salary;

END;


CLOSE emp_cursor;

DEALLOCATE emp_cursor;
```

This cursor updates the salary for each employee in department 1 by 10%.

---

10. What is the performance impact of using Cursors in SQL?

While cursors are useful for row-by-row processing, they can have a **negative impact on performance** because:

1. They require memory and system resources to manage the result set.

2. They can slow down performance as they process each row individually, rather than working with the entire result set at once.

3. They are generally slower than set-based operations (such as UPDATE, INSERT, or DELETE statements that affect multiple rows at once).

It is recommended to use cursors only when necessary, and to explore set-based approaches where possible for better performance.

These answers provide a comprehensive understanding of **cursors in SQL**, including their types, operations, performance implications, and when to use them. Cursors are useful for specific scenarios where row-by-row processing is needed, but they should be used carefully to avoid performance issues.

CRUD OPERATIONS

1. What are CRUD Operations in SQL?

CRUD stands for **Create, Read, Update, and Delete**. These are the four basic operations for managing data in a database:

- **Create**: Inserts new records into a table.

- **Read**: Retrieves or reads data from the table.

- **Update**: Modifies existing data in the table.

- **Delete**: Removes records from the table.

These operations are typically mapped to SQL commands:

- **Create** → INSERT

- **Read** → SELECT

- **Update** → UPDATE

- **Delete** → DELETE

---

2. What is the SQL INSERT Statement (Create operation)?

The INSERT statement is used to add new records (rows) into a table.

Syntax:

sql

Copy code

INSERT INTO table_name (column1, column2, ...)

VALUES (value1, value2, ...);

Example:

sql

Copy code

```
INSERT INTO Employee (emp_id, emp_name, salary, dept_id)

VALUES (1, 'John Doe', 50000, 101);
```

This statement inserts a new row into the Employee table with the given values.

---

3. What is the SQL SELECT Statement (Read operation)?

The SELECT statement is used to retrieve data from a table. You can use various clauses like WHERE, ORDER BY, GROUP BY, etc., to filter, sort, and group the results.

Syntax:

sql

Copy code

```
SELECT column1, column2, ...

FROM table_name

WHERE condition;
```

Example:

sql

Copy code

```
SELECT emp_name, salary FROM Employee WHERE dept_id = 101;
```

This retrieves the names and salaries of employees who belong to department 101.

---

4. What is the SQL UPDATE Statement (Update operation)?

The UPDATE statement is used to modify existing records in a table.

Syntax:

sql

Copy code

UPDATE table_name

SET column1 = value1, column2 = value2, ...

WHERE condition;

Example:

sql

Copy code

UPDATE Employee

SET salary = 55000

WHERE emp_id = 1;

This updates the salary of the employee with emp_id 1 to 55000.

---

5. What is the SQL DELETE Statement (Delete operation)?

The DELETE statement is used to remove records from a table.

Syntax:

sql

Copy code

DELETE FROM table_name

WHERE condition;

Example:

sql

Copy code

DELETE FROM Employee WHERE emp_id = 1;

This deletes the record of the employee with emp_id 1.

6. What are the differences between DELETE, TRUNCATE, and DROP?

- **DELETE**:
    - Removes specific rows from a table based on a condition.
    - Can be rolled back (if inside a transaction).
    - Slower than TRUNCATE because it logs individual row deletions.
    - Does not remove table structure, only data.

- **TRUNCATE**:
    - Removes all rows from a table.
    - Cannot be rolled back (in most RDBMS systems).
    - Faster than DELETE as it doesn't log individual row deletions.
    - Resets identity column values (if any).

- **DROP**:
    - Completely removes the table or other database objects (like views, indexes).
    - Cannot be rolled back.
    - Removes the table structure, not just the data.

---

7. Can we perform multiple CRUD operations in a single SQL query?

Yes, you can perform multiple CRUD operations in a single SQL script, but they need to be written as separate statements. You can't combine them into a single INSERT, SELECT, UPDATE, or DELETE statement.

Example:

sql

Copy code

BEGIN;

INSERT INTO Employee (emp_id, emp_name, salary, dept_id) VALUES (2, 'Jane Doe', 60000, 102);

UPDATE Employee SET salary = 65000 WHERE emp_id = 2;

DELETE FROM Employee WHERE emp_id = 1;

COMMIT;

In this example, we perform an INSERT, UPDATE, and DELETE operation in a single transaction.

---

8. What is a Transaction in SQL, and how does it relate to CRUD operations?

A **transaction** is a sequence of one or more SQL operations that are executed as a single unit. A transaction ensures that either all operations are executed successfully (commit) or, if there is an error, none of the operations are applied (rollback). CRUD operations are often part of a transaction to ensure data integrity.

Syntax:

sql

Copy code

BEGIN TRANSACTION;

-- CRUD operations here

COMMIT; -- Apply changes

ROLLBACK; -- Undo changes in case of an error

Example:

sql

Copy code

```
BEGIN TRANSACTION;

INSERT INTO Employee (emp_id, emp_name, salary, dept_id) VALUES (3,
'Alice', 70000, 103);

UPDATE Employee SET salary = 72000 WHERE emp_id = 3;

COMMIT;
```

This ensures that both the INSERT and UPDATE happen together. If there's an issue, you can ROLLBACK.

---

9. Can you use subqueries within CRUD operations?

Yes, subqueries can be used within CRUD operations. For example, you can use a subquery in the WHERE clause to specify conditions based on the results of another query.

Example:

sql

Copy code

```
UPDATE Employee

SET salary = salary + 5000

WHERE dept_id = (SELECT dept_id FROM Department WHERE dept_name
= 'HR');
```

This updates the salary of all employees in the "HR" department by adding 5000.

---

10. What is the significance of using WHERE in UPDATE and DELETE operations?

The WHERE clause in UPDATE and DELETE operations specifies which rows should be updated or deleted. Without the WHERE clause, all rows in the table would be affected, which could lead to unintended results.

Example:

sql

Copy code

```
-- Without WHERE clause, this will delete all records:

DELETE FROM Employee;


-- With WHERE clause, this deletes only the specified record:

DELETE FROM Employee WHERE emp_id = 2;
```

Always ensure to use a WHERE clause to avoid accidentally deleting or updating all data in the table.

Subqueries

1. What is a Subquery in SQL?

A **subquery** (also called an inner query or nested query) is a query placed inside another query. It is used to perform operations that need to be executed as part of a larger query. Subqueries can return a single value, a list of values, or a table, depending on the type of subquery.

Subqueries are typically used in the SELECT, INSERT, UPDATE, and DELETE statements to extract data that will be used by the outer query.

2. What are the types of Subqueries?

Subqueries can be categorized into the following types:

1. **Single-row Subquery**: Returns a single row and column.

2. **Multi-row Subquery**: Returns multiple rows and one column.

3. **Multi-column Subquery**: Returns multiple rows and columns.

4. **Correlated Subquery**: A subquery that references columns from the outer query.

5. **Non-correlated Subquery**: A subquery that does not reference columns from the outer query.

---

3. What is a Single-row Subquery?

A **single-row subquery** returns only one row of results, usually with a single column. It is often used in the WHERE clause with comparison operators like =, <, >, etc.

Example:

sql

Copy code

SELECT emp_name

FROM Employee

WHERE salary = (SELECT MAX(salary) FROM Employee);

This query returns the name of the employee with the highest salary.

---

4. What is a Multi-row Subquery?

A **multi-row subquery** returns multiple rows and a single column. It is used with operators like IN, ANY, or ALL.

Example:

sql

Copy code

SELECT emp_name

FROM Employee

WHERE dept_id IN (SELECT dept_id FROM Department WHERE location = 'New York');

This query retrieves the names of employees working in departments located in "New York".

---

5. What is a Multi-column Subquery?

A **multi-column subquery** returns multiple rows and multiple columns. It is often used in the IN or EXISTS clause for more complex comparisons.

Example:

sql

Copy code

```
SELECT emp_name

FROM Employee

WHERE (salary, dept_id) IN (SELECT salary, dept_id FROM Employee WHERE salary > 50000);
```

This query retrieves the names of employees whose salary and department match those of employees with a salary greater than 50000.

---

6. What is a Correlated Subquery?

A **correlated subquery** is a subquery that references columns from the outer query. Unlike a non-correlated subquery, which can be executed independently, a correlated subquery depends on the outer query for its values and is executed repeatedly for each row processed by the outer query.

Example:

sql

Copy code

SELECT emp_name, salary

FROM Employee e

WHERE salary > (SELECT AVG(salary) FROM Employee WHERE dept_id = e.dept_id);

In this example, the subquery calculates the average salary for each department, and the outer query retrieves employees whose salary is higher than the average salary for their respective departments.

---

7. What is a Non-correlated Subquery?

A **non-correlated subquery** is independent of the outer query. It does not reference any columns from the outer query. This type of subquery is executed once, and its result is used by the outer query.

Example:

sql

Copy code

SELECT emp_name

FROM Employee

WHERE dept_id = (SELECT dept_id FROM Department WHERE dept_name = 'HR');

Here, the subquery is independent of the outer query and is executed first to find the dept_id of the 'HR' department. The outer query uses this value to find employees in that department.

---

8. How do you use Subqueries in the FROM Clause?

A subquery in the FROM clause can be used to create a temporary table that can be referenced by the outer query. This type of subquery is also called an **inline view**.

Example:

sql

Copy code

SELECT avg_salary, dept_name

FROM (SELECT AVG(salary) AS avg_salary, dept_id FROM Employee GROUP BY dept_id) AS temp

JOIN Department d ON temp.dept_id = d.dept_id;

In this example, the subquery calculates the average salary per department, and the outer query joins it with the Department table to display the average salary along with the department name.

---

9. What are the performance considerations when using Subqueries?

- **Single-row Subqueries**: These tend to perform well, especially when they are used with simple comparison operators. However, performance can degrade if the subquery is complex or if it needs to be executed multiple times (e.g., in a correlated subquery).

- **Multi-row and Multi-column Subqueries**: These can be less efficient because the database may need to handle a larger amount of data. Using operators like IN, ANY, and ALL with these subqueries can slow down performance if the subquery returns a large result set.

- **Correlated Subqueries**: These are generally less efficient because the subquery is executed repeatedly for each row of the outer query, which can lead to performance bottlenecks.

- **Indexes**: Ensuring that columns used in subqueries, especially in WHERE clauses, are indexed can improve performance.

---

10. Can you use Subqueries in UPDATE, DELETE, and INSERT statements?

Yes, subqueries can be used in UPDATE, DELETE, and INSERT statements for more complex operations. Here's how:

- **UPDATE**:

sql

Copy code

UPDATE Employee

SET salary = salary + 5000

WHERE dept_id = (SELECT dept_id FROM Department WHERE dept_name = 'HR');

- **DELETE**:

sql

Copy code

DELETE FROM Employee

WHERE dept_id = (SELECT dept_id FROM Department WHERE dept_name = 'HR');

- **INSERT**:

sql

Copy code

INSERT INTO Employee (emp_id, emp_name, salary, dept_id)

SELECT emp_id, emp_name, salary, dept_id

FROM TempEmployee

WHERE hire_date > '2023-01-01';

In these cases, subqueries help to identify which rows should be updated, deleted, or inserted based on data from other tables or conditions.

DML

1. What is DML (Data Manipulation Language)?

DML refers to a subset of SQL commands that deal with the manipulation of data in database tables. DML operations allow users to **retrieve**, **insert**, **update**, and **delete** data stored in the database. These operations work with the data stored in the tables, rather than the database schema.

The primary DML operations are:

- **SELECT**: Retrieves data from the database.

- **INSERT**: Adds new rows of data to a table.

- **UPDATE**: Modifies existing data in a table.

- **DELETE**: Removes data from a table.

---

2. What is the INSERT Statement in DML?

The INSERT statement is used to add new rows into a table.

Syntax:

sql

Copy code

INSERT INTO table_name (column1, column2, ...)

VALUES (value1, value2, ...);

Example:

sql

Copy code

INSERT INTO Employee (emp_id, emp_name, salary, dept_id)

VALUES (1, 'John Doe', 50000, 101);

This query inserts a new employee into the Employee table with the provided values for emp_id, emp_name, salary, and dept_id.

---

3. What is the SELECT Statement in DML?

The SELECT statement is used to retrieve data from one or more tables. It can be used with various clauses like WHERE, ORDER BY, GROUP BY, and more to filter, sort, and aggregate data.

Syntax:

sql

Copy code

SELECT column1, column2, ...

FROM table_name

WHERE condition;

Example:

sql

Copy code

SELECT emp_name, salary FROM Employee WHERE dept_id = 101;

This query retrieves the names and salaries of employees who belong to department 101.

---

4. What is the UPDATE Statement in DML?

The UPDATE statement is used to modify the existing records in a table. You can update one or more columns in one or more rows based on a specified condition.

Syntax:

sql

Copy code

UPDATE table_name

SET column1 = value1, column2 = value2, ...

WHERE condition;

Example:

sql

Copy code

UPDATE Employee

SET salary = 55000

WHERE emp_id = 1;

This query updates the salary of the employee with emp_id 1 to 55000.

---

5. What is the DELETE Statement in DML?

The DELETE statement is used to remove rows from a table based on a specified condition. It removes data but does not affect the table structure.

Syntax:

sql

Copy code

DELETE FROM table_name

WHERE condition;

Example:

sql

Copy code

DELETE FROM Employee WHERE emp_id = 1;

This query deletes the employee record where the emp_id is 1.

---

6. What is the role of the WHERE clause in DML statements?

The WHERE clause is essential in DML statements like UPDATE and DELETE because it specifies the condition that rows must meet to be

affected. Without the WHERE clause, the operation would affect all rows in the table.

Example:

- **Without WHERE**:

sql

Copy code

DELETE FROM Employee;

This deletes all records from the Employee table.

- **With WHERE**:

sql

Copy code

DELETE FROM Employee WHERE emp_id = 1;

This deletes only the record where emp_id equals 1.

---

7. What is the difference between TRUNCATE and DELETE in DML?

- **DELETE**:
    - Removes specific rows from a table based on a condition.
    - Can be rolled back (if inside a transaction).
    - Slower than TRUNCATE because it logs individual row deletions.
    - Does not reset the table structure or any associated identity columns.

- **TRUNCATE**:
    - Removes all rows from a table.
    - Cannot be rolled back (in most databases).

- o Faster than DELETE because it does not log individual row deletions.
- o Resets identity columns (if any).

---

8. Can you perform multiple DML operations in a single transaction?

Yes, multiple DML operations can be executed as a single **transaction**. Transactions ensure that either all changes are applied (commit) or none (rollback) in case of an error.

Example:

sql

Copy code

BEGIN;

INSERT INTO Employee (emp_id, emp_name, salary, dept_id) VALUES (2, 'Jane Doe', 60000, 102);

UPDATE Employee SET salary = 65000 WHERE emp_id = 2;

DELETE FROM Employee WHERE emp_id = 1;

COMMIT; -- All changes are saved

In case of an error, a ROLLBACK can be used to undo all changes made during the transaction.

---

9. What is the significance of COMMIT and ROLLBACK in DML operations?

- **COMMIT**:
  - o It permanently saves the changes made during a transaction.
  - o Once a COMMIT is executed, the changes cannot be undone.

- **ROLLBACK**:
    - It undoes all changes made during the current transaction.
    - If an error occurs or if the user wants to discard changes, ROLLBACK is used.

Example:

sql

Copy code

BEGIN;

UPDATE Employee SET salary = 70000 WHERE emp_id = 2;

-- If no errors occur

COMMIT;

---

10. What are the potential risks of using DML operations in a production database?

When using DML operations in a production environment, there are several risks to be aware of:

- **Accidental Data Loss**: DELETE or UPDATE without the proper WHERE condition could result in unintended data loss.
- **Data Integrity Issues**: If transactions are not properly managed, it could lead to inconsistent or incorrect data in the database.
- **Performance Degradation**: Frequent updates or deletions in large tables can degrade performance.
- **Concurrency Issues**: Multiple users performing DML operations simultaneously can cause data conflicts, which is why transactions are used to ensure data consistency.

To mitigate these risks, it's crucial to:

- Always use the WHERE clause to limit the scope of updates and deletions.

- Use transactions to ensure data consistency and rollback in case of errors.

- Test queries in a non-production environment before applying them to the live database.

DDL

1. What is DDL (Data Definition Language)?

DDL refers to a subset of SQL commands used to define and modify the structure of database objects such as tables, indexes, and schemas. DDL operations do not manipulate data directly but instead deal with database schema, which is the organization of the data.

The main DDL operations are:

- **CREATE**: Defines a new database object (e.g., table, view, index).

- **ALTER**: Modifies an existing database object.

- **DROP**: Deletes an existing database object.

- **TRUNCATE**: Removes all rows from a table but does not affect its structure.

---

2. What is the CREATE Statement in DDL?

The CREATE statement is used to define a new database object such as a table, view, index, or schema.

Syntax:

sql

Copy code

```sql
CREATE TABLE table_name (

    column1 datatype,

    column2 datatype,

    ...

);
```

Example:

sql

Copy code

```sql
CREATE TABLE Employee (

    emp_id INT PRIMARY KEY,

    emp_name VARCHAR(50),

    salary DECIMAL(10, 2),

    dept_id INT

);
```

This query creates a new table Employee with columns for emp_id, emp_name, salary, and dept_id.

---

3. What is the ALTER Statement in DDL?

The ALTER statement is used to modify an existing database object, such as adding or deleting columns from a table, changing a column's datatype, or modifying constraints.

Syntax:

- **Add a column**:

sql

Copy code

```sql
ALTER TABLE table_name ADD column_name datatype;
```

- **Modify a column**:

sql

Copy code

ALTER TABLE table_name MODIFY column_name datatype;

- **Drop a column**:

sql

Copy code

ALTER TABLE table_name DROP COLUMN column_name;

Example:

sql

Copy code

ALTER TABLE Employee ADD email VARCHAR(100);

This query adds a new column email to the Employee table.

---

4. What is the DROP Statement in DDL?

The DROP statement is used to delete an existing database object, such as a table, index, or view. Once an object is dropped, it cannot be recovered unless a backup is available.

Syntax:

sql

Copy code

DROP TABLE table_name;

Example:

sql

Copy code

DROP TABLE Employee;

This query deletes the Employee table from the database permanently.

---

5. What is the TRUNCATE Statement in DDL?

The TRUNCATE statement is used to remove all rows from a table, but unlike DELETE, it does not log individual row deletions and cannot be rolled back (in most cases).

Syntax:

sql

Copy code

TRUNCATE TABLE table_name;

Example:

sql

Copy code

TRUNCATE TABLE Employee;

This query removes all rows from the Employee table but leaves the structure of the table intact.

---

6. What is the difference between DROP and TRUNCATE?

- **DROP**:
  - Completely removes the table or other database objects (like views, indexes).
  - Cannot be rolled back.
  - Removes both the data and the structure (table, columns, constraints).
  - Frees the space occupied by the table.
- **TRUNCATE**:

- o Removes all rows from the table but retains the structure.

- o Can be rolled back (in some DBMS like MySQL and SQL Server).

- o Does not fire triggers.

- o Does not reset the identity column (unless specifically configured).

---

7. What are Constraints in DDL and how do they work?

Constraints are rules enforced on columns to maintain data integrity. They are defined when creating or altering tables.

Types of constraints include:

1. **PRIMARY KEY**: Uniquely identifies each row in a table.

2. **FOREIGN KEY**: Ensures referential integrity by linking a column to another table's primary key.

3. **UNIQUE**: Ensures all values in a column are distinct.

4. **CHECK**: Ensures that all values in a column satisfy a specific condition.

5. **NOT NULL**: Ensures that a column cannot contain NULL values.

Example:

sql

Copy code

```sql
CREATE TABLE Employee (
    emp_id INT PRIMARY KEY,
    emp_name VARCHAR(50) NOT NULL,
    salary DECIMAL(10, 2),
    dept_id INT,
```

FOREIGN KEY (dept_id) REFERENCES Department(dept_id)

);

In this example, the emp_id is the primary key, emp_name cannot be NULL, and dept_id is a foreign key that references the Department table.

---

8. What is the CREATE INDEX Statement in DDL?

The CREATE INDEX statement is used to create an index on a table. Indexes improve the speed of data retrieval operations, but they can slow down insert, update, and delete operations because the index also needs to be updated.

Syntax:

sql

Copy code

```
CREATE INDEX index_name

ON table_name (column_name);
```

Example:

sql

Copy code

```
CREATE INDEX idx_salary

ON Employee (salary);
```

This query creates an index named idx_salary on the salary column of the Employee table to speed up queries involving salary.

---

9. What is the CREATE VIEW Statement in DDL?

The CREATE VIEW statement is used to create a virtual table (view) based on the result of a SELECT query. Views do not store data themselves but provide a way to query data in a simplified manner.

Syntax:

sql

Copy code

CREATE VIEW view_name AS

SELECT column1, column2, ...

FROM table_name

WHERE condition;

Example:

sql

Copy code

CREATE VIEW HighSalaryEmployees AS

SELECT emp_name, salary

FROM Employee

WHERE salary > 60000;

This query creates a view HighSalaryEmployees that shows the names and salaries of employees earning more than 60,000.

---

10. What are the performance considerations when using DDL operations?

- **CREATE**: Creating tables or indexes can be resource-intensive, especially with large datasets. It's important to optimize the schema design to ensure that operations remain efficient.

- **ALTER**: Altering table structures (e.g., adding columns or changing data types) can result in significant performance overhead, particularly with large tables. Ensure that such operations are planned and executed during off-peak hours when possible.

- **DROP**: Dropping tables or indexes is relatively fast but should be used with caution since it permanently removes data and schema.

- **TRUNCATE**: While TRUNCATE is faster than DELETE, it can lead to performance issues if used frequently, especially if the table is large and has many foreign key relationships.

Best practices to mitigate performance issues include:

- Index optimization.

- Performing schema changes during low-traffic periods.

- Regularly reviewing and optimizing the table structure to prevent unnecessary complexity.